

introduction-to-numpy

October 16, 2024

1 Intro To Numpy!

Helpful Jupyter Tip! : You can view the documentation of a function by moving your cursor on top of it and pressing **shift+tab**

1.1 Why NumPy?

You can do numerical calculations using pure Python. In the beginning, you might think Python is fast but once your data gets large, you'll start to notice slow downs.

One of the main reasons you use NumPy is because it's fast. Behind the scenes, the code has been optimized to run using C. Which is another programming language, which can do things much faster than Python.

The benefit of this being behind the scenes is you don't need to know any C to take advantage of it. You can write your numerical computations in Python using NumPy and get the added speed benefits.

If your curious as to what causes this speed benefit, it's a process called vectorization. [Vectorization](#) aims to do calculations by avoiding loops as loops can create potential bottlenecks.

NumPy achieves vectorization through a process called [broadcasting](#).

```
[2]: import numpy as np

print(np.__version__)
```

1.26.4

1.2 1. DataTypes and attributes

Note: Important to remember the main type in NumPy is `ndarray`, even seemingly different kinds of arrays are still `ndarray`'s. This means an operation you do on one array, will work on another.

```
[3]: # 1-dimensional array, also referred to as a vector
a1 = np.array([1, 2, 3])

# 2-dimensional array, also referred to as matrix
a2 = np.array([[1, 2.0, 3.3],
               [4, 5, 6.5]])
```

```
# 3-dimensional array, also referred to as a matrix
a3 = np.array([[[1, 2, 3],
                [4, 5, 6],
                [7, 8, 9]],
               [[10, 11, 12],
                [13, 14, 15],
                [16, 17, 18]]])
```

```
[4]: a1.shape, a1.ndim, a1.dtype, a1.size, type(a1)
```

```
[4]: ((3,), 1, dtype('int64'), 3, numpy.ndarray)
```

```
[5]: a2.shape, a2.ndim, a2.dtype, a2.size, type(a2)
```

```
[5]: ((2, 3), 2, dtype('float64'), 6, numpy.ndarray)
```

```
[6]: a3.shape, a3.ndim, a3.dtype, a3.size, type(a3)
```

```
[6]: ((2, 3, 3), 3, dtype('int64'), 18, numpy.ndarray)
```

```
[7]: a1
```

```
[7]: array([1, 2, 3])
```

```
[8]: a2
```

```
[8]: array([[1. , 2. , 3.3],
           [4. , 5. , 6.5]])
```

```
[9]: a3
```

```
[9]: array([[[ 1,  2,  3],
              [ 4,  5,  6],
              [ 7,  8,  9]],
            [[10, 11, 12],
              [13, 14, 15],
              [16, 17, 18]]])
```

1.3 2. Creating arrays

- `np.array()`
- `np.ones()`
- `np.zeros()`
- `np.random.rand(5, 3)`
- `np.random.randint(10, size=5)`
- `np.random.seed()` - pseudo random numbers
- `np.arange()`

```
[5]: # Create a simple array
```

```
[6]: # Use the array funcs and list details of this array
```

```
[7]: # Create an array of ones
```

```
[11]: # cast the created array into another dtype using .astype("dtype")
```

```
[12]: # Create an array of zeros
```

```
[13]: # Create an array within a range of values
```

```
[16]: # Random array
```

```
[17]: # Random array of floats (between 0 & 1)
```

NumPy uses pseudo-random numbers, which means, the numbers look random but aren't really, they're predetermined.

For consistency, you might want to keep the random numbers you generate similar throughout experiments.

To do this, you can use `np.random.seed()`.

What this does is it tells NumPy, "Hey, I want you to create random numbers but keep them aligned with the seed."

Let's see it.

```
[18]: # Set random seed to 0
```

With `np.random.seed()` set, every time you run the cell above, the same random numbers will be generated.

What if `np.random.seed()` wasn't set?

Every time you run the cell below, a new set of numbers will appear.

```
[27]: # Make more random numbers
np.random.randint(10, size=(5, 3))
```

```
[27]: array([[6, 7, 7],
            [8, 1, 5],
            [9, 8, 9],
            [4, 3, 0],
            [3, 5, 0]])
```

Let's see it in action again, we'll stay consistent and set the random seed to 0.

```
[28]: # Set random seed to same number as above
np.random.seed(0)
```

```
# The same random numbers come out
np.random.randint(10, size=(5, 3))
```

```
[28]: array([[5, 0, 3],
            [3, 7, 9],
            [3, 5, 2],
            [4, 7, 6],
            [8, 8, 1]])
```

Because `np.random.seed()` is set to 0, the random numbers are the same as the cell with `np.random.seed()` set to 0 as well.

Setting `np.random.seed()` is not 100% necessary but it's helpful to keep numbers the same throughout your experiments.

For example, say you wanted to split your data randomly into training and test sets.

Every time you randomly split, you might get different rows in each set.

If you shared your work with someone else, they'd get different rows in each set too.

Setting `np.random.seed()` ensures there's still randomness, it just makes the randomness repeatable. Hence the 'pseudo-random' numbers.

1.3.1 What unique values are in the array `a3`?

Now you've seen a few different ways to create arrays, as an exercise, try find out what NumPy function you could use to find the unique values are within the `a3` array.

You might want to search some like, "how to find the unique values in a numpy array".

```
[30]: # Your code here
```

1.4 3. Viewing arrays and matrices (indexing)

Remember, because arrays and matrices are both `ndarray`'s, they can be viewed in similar ways.

Let's check out our 3 arrays again.

```
[31]: a1
```

```
[31]: array([1, 2, 3])
```

```
[32]: a2
```

```
[32]: array([[1. , 2. , 3.3],
            [4. , 5. , 6.5]])
```

```
[33]: a3
```

```
[33]: array([[ 1,  2,  3],
            [ 4,  5,  6],
```

```

        [ 7,  8,  9]],

        [[10, 11, 12],
         [13, 14, 15],
         [16, 17, 18]]])

```

Array shapes are always listed in the format (row, column, n, n, n...) where n is optional extra dimensions.

```
[34]: a1[0]
```

```
[34]: np.int64(1)
```

```
[35]: a2[0]
```

```
[35]: array([1. , 2. , 3.3])
```

```
[36]: a3[0]
```

```
[36]: array([[1, 2, 3],
            [4, 5, 6],
            [7, 8, 9]])
```

```
[37]: # Get 2nd row (index 1) of a2
      a2[1]
```

```
[37]: array([4. , 5. , 6.5])
```

```
[38]: # Get the first 2 values of the first 2 rows of both arrays
      a3[:2, :2, :2]
```

```
[38]: array([[[ 1,  2],
              [ 4,  5]],

            [[10, 11],
             [13, 14]]])
```

This takes a bit of practice, especially when the dimensions get higher. Usually, it takes me a little trial and error of trying to get certain values, viewing the output in the notebook and trying again.

NumPy arrays get printed from outside to inside. This means the number at the end of the shape comes first, and the number at the start of the shape comes last.

```
[39]: a4 = np.random.randint(10, size=(2, 3, 4, 5))
      a4
```

```
[39]: array([[[[6, 7, 7, 8, 1],
              [5, 9, 8, 9, 4],
              [3, 0, 3, 5, 0],
```

```

[2, 3, 8, 1, 3]],

[[3, 3, 7, 0, 1],
 [9, 9, 0, 4, 7],
 [3, 2, 7, 2, 0],
 [0, 4, 5, 5, 6]],

[[8, 4, 1, 4, 9],
 [8, 1, 1, 7, 9],
 [9, 3, 6, 7, 2],
 [0, 3, 5, 9, 4]]],

[[[4, 6, 4, 4, 3],
  [4, 4, 8, 4, 3],
  [7, 5, 5, 0, 1],
  [5, 9, 3, 0, 5]],

 [[0, 1, 2, 4, 2],
  [0, 3, 2, 0, 7],
  [5, 9, 0, 2, 7],
  [2, 9, 2, 3, 3]],

 [[2, 3, 4, 1, 2],
  [9, 1, 4, 6, 8],
  [2, 3, 0, 0, 6],
  [0, 6, 3, 3, 8]]]])

```

```
[40]: a4.shape
```

```
[40]: (2, 3, 4, 5)
```

```
[41]: # Get only the first 4 numbers of each single vector
a4[:, :, :, :4]
```

```
[41]: array([[[[6, 7, 7, 8],
               [5, 9, 8, 9],
               [3, 0, 3, 5],
               [2, 3, 8, 1]],

               [[3, 3, 7, 0],
               [9, 9, 0, 4],
               [3, 2, 7, 2],
               [0, 4, 5, 5]],

               [[8, 4, 1, 4],
               [8, 1, 1, 7],
```

```

[9, 3, 6, 7],
[0, 3, 5, 9]]],

[[[4, 6, 4, 4],
  [4, 4, 8, 4],
  [7, 5, 5, 0],
  [5, 9, 3, 0]],

 [[0, 1, 2, 4],
  [0, 3, 2, 0],
  [5, 9, 0, 2],
  [2, 9, 2, 3]],

 [[2, 3, 4, 1],
  [9, 1, 4, 6],
  [2, 3, 0, 0],
  [0, 6, 3, 3]]]])

```

`a4`'s shape is (2, 3, 4, 5), this means it gets displayed like so: * Inner most array = size 5 * Next array = size 4 * Next array = size 3 * Outer most array = size 2

1.5 4. Manipulating and comparing arrays

- Arithmetic
 - +, -, *, /, //, **, %
 - `np.exp()`
 - `np.log()`
 - [Dot product](#) - `np.dot()`
 - Broadcasting
- Aggregation
 - `np.sum()` - faster than Python's `.sum()` for NumPy arrays
 - `np.mean()`
 - `np.std()`
 - `np.var()`
 - `np.min()`
 - `np.max()`
 - `np.argmax()` - find index of minimum value
 - `np.argmin()` - find index of maximum value
 - These work on all `ndarray`'s
 - * `a4.min(axis=0)` – you can use axis as well
- Reshaping
 - `np.reshape()`
- Transposing
 - `a3.T`
- Comparison operators
 - >
 - <

```
- <=
- >=
- x != 3
- x == 3
- np.sum(x > 3)
```

1.5.1 Arithmetic

```
[42]: a1
```

```
[42]: array([1, 2, 3])
```

```
[43]: ones = np.ones(3)
ones
```

```
[43]: array([1., 1., 1.])
```

```
[44]: # Add two arrays
a1 + ones
```

```
[44]: array([2., 3., 4.])
```

```
[45]: # Subtract two arrays
a1 - ones
```

```
[45]: array([0., 1., 2.])
```

```
[46]: # Multiply two arrays
a1 * ones
```

```
[46]: array([1., 2., 3.])
```

```
[47]: # Multiply two arrays
a1 * a2
```

```
[47]: array([[ 1. ,  4. ,  9.9],
           [ 4. , 10. , 19.5]])
```

```
[48]: a1.shape, a2.shape
```

```
[48]: ((3,), (2, 3))
```

```
[49]: # This will error as the arrays have a different number of dimensions (2, 3) vs.
      ↪ (2, 3, 3)
a2 * a3
```

```
-----
ValueError
```

```
Traceback (most recent call last)
```

```
Cell In[49], line 2
```



```
1 # This will error as the arrays have a different number of dimensions
↪(2, 3) vs. (2, 3, 3)
-----> 2 a2 * a3
```

```
ValueError: operands could not be broadcast together with shapes (2,3) (2,3,3)
```

```
[50]: a3
```

```
[50]: array([[[ 1,  2,  3],
             [ 4,  5,  6],
             [ 7,  8,  9]],

           [[10, 11, 12],
            [13, 14, 15],
            [16, 17, 18]]])
```

1.5.2 Broadcasting

- What is broadcasting?
 - Broadcasting is a feature of NumPy which performs an operation across multiple dimensions of data without replicating the data. This saves time and space. For example, if you have a 3x3 array (A) and want to add a 1x3 array (B), NumPy will add the row of (B) to every row of (A).
- Rules of Broadcasting
 1. If the two arrays differ in their number of dimensions, the shape of the one with fewer dimensions is padded with ones on its leading (left) side.
 2. If the shape of the two arrays does not match in any dimension, the array with shape equal to 1 in that dimension is stretched to match the other shape.
 3. If in any dimension the sizes disagree and neither is equal to 1, an error is raised.

The broadcasting rule: In order to broadcast, the size of the trailing axes for both arrays in an operation must be either the same size or one of them must be one.

```
[51]: a1
```

```
[51]: array([1, 2, 3])
```

```
[52]: a1.shape
```

```
[52]: (3,)
```

```
[53]: a2.shape
```

```
[53]: (2, 3)
```

```
[54]: a2
```

```
[54]: array([[1. , 2. , 3.3],
           [4. , 5. , 6.5]])
```

```
[55]: a1 + a2
```

```
[55]: array([[2. , 4. , 6.3],
           [5. , 7. , 9.5]])
```

```
[56]: a2 + 2
```

```
[56]: array([[3. , 4. , 5.3],
           [6. , 7. , 8.5]])
```

```
[57]: # Raises an error because there's a shape mismatch (2, 3) vs. (2, 3, 3)
      a2 + a3
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[57], line 2
      1 # Raises an error because there's a shape mismatch (2, 3) vs. (2, 3, 3)
----> 2 a2 + a3

ValueError: operands could not be broadcast together with shapes (2,3) (2,3,3)
```

```
[58]: # Divide two arrays
      a1 / ones
```

```
[58]: array([1., 2., 3.])
```

```
[59]: # Divide using floor division
      a2 // a1
```

```
[59]: array([[1., 1., 1.],
           [4., 2., 2.]])
```

```
[60]: # Take an array to a power
      a1 ** 2
```

```
[60]: array([1, 4, 9])
```

```
[61]: # You can also use np.square()
      np.square(a1)
```

```
[61]: array([1, 4, 9])
```

```
[62]: # Modulus divide (what's the remainder)
      a1 % 2
```

```
[62]: array([1, 0, 1])
```

You can also find the log or exponential of an array using `np.log()` and `np.exp()`.

```
[63]: # Find the log of an array  
np.log(a1)
```

```
[63]: array([0.          , 0.69314718, 1.09861229])
```

```
[64]: # Find the exponential of an array  
np.exp(a1)
```

```
[64]: array([ 2.71828183,  7.3890561 , 20.08553692])
```

1.5.3 Aggregation

Aggregation - bringing things together, doing a similar thing on a number of things.

```
[65]: sum(a1)
```

```
[65]: np.int64(6)
```

```
[66]: np.sum(a1)
```

```
[66]: np.int64(6)
```

Tip: Use NumPy's `np.sum()` on NumPy arrays and Python's `sum()` on Python lists.

```
[67]: massive_array = np.random.random(100000)  
massive_array.size, type(massive_array)
```

```
[67]: (100000, numpy.ndarray)
```

```
[68]: %timeit sum(massive_array) # Python sum()  
%timeit np.sum(massive_array) # NumPy np.sum()
```

3.93 ms \pm 145 s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

20.5 s \pm 698 ns per loop (mean \pm std. dev. of 7 runs, 10,000 loops each)

Notice `np.sum()` is faster on the Numpy array (`numpy.ndarray`) than Python's `sum()`.

Now let's try it out on a Python list.

```
[69]: import random  
massive_list = [random.randint(0, 10) for i in range(100000)]  
len(massive_list), type(massive_list)
```

```
[69]: (100000, list)
```

```
[70]: massive_list[:10]
```

```
[70]: [8, 9, 1, 0, 0, 6, 2, 8, 6, 3]
```

```
[71]: %timeit sum(massive_list)
      %timeit np.sum(massive_list)
```

419 s \pm 6.74 s per loop (mean \pm std. dev. of 7 runs, 1,000 loops each)

2.72 ms \pm 118 s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

NumPy's np.sum() is still fast but Python's sum() is faster on Python lists.

```
[72]: a2
```

```
[72]: array([[1. , 2. , 3.3],
           [4. , 5. , 6.5]])
```

```
[73]: # Find the mean
      np.mean(a2)
```

```
[73]: np.float64(3.6333333333333333)
```

```
[74]: # Find the max
      np.max(a2)
```

```
[74]: np.float64(6.5)
```

```
[75]: # Find the min
      np.min(a2)
```

```
[75]: np.float64(1.0)
```

```
[76]: # Find the standard deviation
      np.std(a2)
```

```
[76]: np.float64(1.8226964152656422)
```

```
[77]: # Find the variance
      np.var(a2)
```

```
[77]: np.float64(3.3222222222222224)
```

```
[78]: # The standard deviation is the square root of the variance
      np.sqrt(np.var(a2))
```

```
[78]: np.float64(1.8226964152656422)
```

What's mean?

Mean is the same as average. You can find the average of a set of numbers by adding them up and dividing them by how many there are.

What's standard deviation?

Standard deviation is a measure of how spread out numbers are.

What's variance?

The **variance** is the averaged squared differences of the mean.

To work it out, you: 1. Work out the mean 2. For each number, subtract the mean and square the result 3. Find the average of the squared differences

```
[79]: # Demo of variance
high_var_array = np.array([1, 100, 200, 300, 4000, 5000])
low_var_array = np.array([2, 4, 6, 8, 10])

np.var(high_var_array), np.var(low_var_array)
```

```
[79]: (np.float64(4296133.472222221), np.float64(8.0))
```


```
[80]: np.std(high_var_array), np.std(low_var_array)
```

```
[80]: (np.float64(2072.711623024829), np.float64(2.8284271247461903))
```


```
[81]: # The standard deviation is the square root of the variance
np.sqrt(np.var(high_var_array))
```

```
[81]: np.float64(2072.711623024829)
```

```
[82]: %matplotlib inline
import matplotlib.pyplot as plt
plt.hist(high_var_array)
plt.show()
```



```
[83]: plt.hist(low_var_array)  
plt.show()
```



1.5.4 Reshaping

```
[84]: a2
```

```
[84]: array([[1. , 2. , 3.3],  
          [4. , 5. , 6.5]])
```

```
[85]: a2.shape
```

```
[85]: (2, 3)
```

```
[86]: a2 + a3
```

```
-----  
ValueError                                Traceback (most recent call last)  
Cell In[86], line 1  
----> 1 a2 + a3  
  
ValueError: operands could not be broadcast together with shapes (2,3) (2,3,3)
```

```
[ ]: a2.reshape(2, 3, 1)
```

```
[87]: a2.reshape(2, 3, 1) + a3
```

```
[87]: array([[[ 2. ,  3. ,  4. ],
             [ 6. ,  7. ,  8. ],
             [10.3, 11.3, 12.3]],

           [[14. , 15. , 16. ],
             [18. , 19. , 20. ],
             [22.5, 23.5, 24.5]]])
```

1.5.5 Transpose

A transpose reverses the order of the axes.

For example, an array with shape (2, 3) becomes (3, 2).

```
[88]: a2.shape
```

```
[88]: (2, 3)
```

```
[89]: a2.T
```

```
[89]: array([[1. , 4. ],
            [2. , 5. ],
            [3.3, 6.5]])
```

```
[90]: a2.transpose()
```

```
[90]: array([[1. , 4. ],
            [2. , 5. ],
            [3.3, 6.5]])
```

```
[91]: a2.T.shape
```

```
[91]: (3, 2)
```

For larger arrays, the default value of a transpose is to swap the first and last axes.

For example, (5, 3, 3) -> (3, 3, 5).

```
[92]: matrix = np.random.random(size=(5, 3, 3))
matrix
```

```
[92]: array([[[0.59816399, 0.17370251, 0.49752936],
             [0.51231935, 0.41529741, 0.44150892],
             [0.96844105, 0.23242417, 0.90336451]],

           [[0.35172075, 0.56481088, 0.57771134],
             [0.73115238, 0.88762934, 0.37368847],
             [0.35104994, 0.11873224, 0.72324236]],
```



```

[[0.93202688, 0.09600718, 0.4330638 ],
 [0.71979707, 0.06689016, 0.20815443],
 [0.55415679, 0.08416165, 0.88953996]],

[[0.00301345, 0.30163886, 0.12337636],
 [0.13435611, 0.51987339, 0.05418991],
 [0.11426417, 0.19005404, 0.61364183]],

[[0.23385887, 0.13555752, 0.32546415],
 [0.81922614, 0.94551446, 0.12975713],
 [0.35431267, 0.37758386, 0.07987885]]])

```

```
[93]: matrix.shape
```

```
[93]: (5, 3, 3)
```

```
[94]: matrix.T
```

```

[94]: array([[0.59816399, 0.35172075, 0.93202688, 0.00301345, 0.23385887],
 [0.51231935, 0.73115238, 0.71979707, 0.13435611, 0.81922614],
 [0.96844105, 0.35104994, 0.55415679, 0.11426417, 0.35431267]],

 [[0.17370251, 0.56481088, 0.09600718, 0.30163886, 0.13555752],
 [0.41529741, 0.88762934, 0.06689016, 0.51987339, 0.94551446],
 [0.23242417, 0.11873224, 0.08416165, 0.19005404, 0.37758386]],

 [[0.49752936, 0.57771134, 0.4330638 , 0.12337636, 0.32546415],
 [0.44150892, 0.37368847, 0.20815443, 0.05418991, 0.12975713],
 [0.90336451, 0.72324236, 0.88953996, 0.61364183, 0.07987885]])

```

```
[95]: matrix.T.shape
```

```
[95]: (3, 3, 5)
```

```

[96]: # Check to see if the reverse shape is same as tranpose shape
matrix.T.shape == matrix.shape[::-1]

```

```
[96]: True
```

```

[97]: # Check to see if the first and last axes are swapped
matrix.T == matrix.swapaxes(0, -1) # swap first (0) and last (-1) axes

```

```

[97]: array([[[ True,  True,  True,  True,  True],
 [ True,  True,  True,  True,  True],
 [ True,  True,  True,  True,  True]],

 [[ True,  True,  True,  True,  True],

```

```
[ True,  True,  True,  True,  True],
 [ True,  True,  True,  True,  True]],

[[ True,  True,  True,  True,  True],
 [ True,  True,  True,  True,  True],
 [ True,  True,  True,  True,  True]])
```

You can see more advanced forms of tranposing in the NumPy documentation under [numpy.transpose](#).

1.5.6 Dot product

The main two rules for dot product to remember are:

1. The **inner dimensions** must match:
 - (3, 2) @ (3, 2) won't work
 - (2, 3) @ (3, 2) will work
 - (3, 2) @ (2, 3) will work
2. The resulting matrix has the shape of the **outer dimensions**:
 - (2, 3) @ (3, 2) -> (2, 2)
 - (3, 2) @ (2, 3) -> (3, 3)

Note: In NumPy, `np.dot()` and `@` can be used to acheive the same result for 1-2 dimension arrays. However, their behaviour begins to differ at arrays with 3+ dimensions.

```
[98]: np.random.seed(0)
      mat1 = np.random.randint(10, size=(3, 3))
      mat2 = np.random.randint(10, size=(3, 2))

      mat1.shape, mat2.shape
```

```
[98]: ((3, 3), (3, 2))
```

```
[99]: mat1
```

```
[99]: array([[5, 0, 3],
             [3, 7, 9],
             [3, 5, 2]])
```

```
[100]: mat2
```

```
[100]: array([[4, 7],
             [6, 8],
             [8, 1]])
```

```
[101]: np.dot(mat1, mat2)
```

```
[101]: array([[ 44,  38],
              [126,  86],
              [ 58,  63]])
```

```
[102]: # Can also achieve np.dot() with "@"
       # (however, they may behave differently at 3D+ arrays)
       mat1 @ mat2
```

```
[102]: array([[ 44,  38],
              [126,  86],
              [ 58,  63]])
```

```
[103]: np.random.seed(0)
       mat3 = np.random.randint(10, size=(4,3))
       mat4 = np.random.randint(10, size=(4,3))
       mat3
```

```
[103]: array([[5, 0, 3],
              [3, 7, 9],
              [3, 5, 2],
              [4, 7, 6]])
```

```
[104]: mat4
```

```
[104]: array([[8, 8, 1],
              [6, 7, 7],
              [8, 1, 5],
              [9, 8, 9]])
```

```
[105]: # This will fail as the inner dimensions of the matrices do not match
       np.dot(mat3, mat4)
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[105], line 2
      1 # This will fail as the inner dimensions of the matrices do not match
----> 2 np.dot(mat3, mat4)

ValueError: shapes (4,3) and (4,3) not aligned: 3 (dim 1) != 4 (dim 0)
```

```
[106]: mat3.T.shape
```

```
[106]: (3, 4)
```

```
[107]: # Dot product
       np.dot(mat3.T, mat4)
```

```
[107]: array([[118, 96, 77],
             [145, 110, 137],
             [148, 137, 130]])
```

```
[108]: # Element-wise multiplication, also known as Hadamard product
mat3 * mat4
```

```
[108]: array([[40, 0, 3],
             [18, 49, 63],
             [24, 5, 10],
             [36, 56, 54]])
```

1.5.7 Dot product practical example, nut butter sales

```
[109]: np.random.seed(0)
sales_amounts = np.random.randint(20, size=(5, 3))
sales_amounts
```

```
[109]: array([[12, 15, 0],
             [ 3,  3,  7],
             [ 9, 19, 18],
             [ 4,  6, 12],
             [ 1,  6,  7]])
```

```
[110]: weekly_sales = pd.DataFrame(sales_amounts,
                                   index=["Mon", "Tues", "Wed", "Thurs", "Fri"],
                                   columns=["Almond butter", "Peanut butter", "Cashew_
↪butter"])
weekly_sales
```

```
[110]:
```

	Almond butter	Peanut butter	Cashew butter
Mon	12	15	0
Tues	3	3	7
Wed	9	19	18
Thurs	4	6	12
Fri	1	6	7

```
[111]: prices = np.array([10, 8, 12])
prices
```

```
[111]: array([10,  8, 12])
```

```
[112]: butter_prices = pd.DataFrame(prices.reshape(1, 3),
                                   index=["Price"],
                                   columns=["Almond butter", "Peanut butter", "Cashew_
↪butter"])
butter_prices.shape
```

```
[112]: (1, 3)
```

```
[113]: weekly_sales.shape
```

```
[113]: (5, 3)
```

```
[114]: # Find the total amount of sales for a whole day
total_sales = prices.dot(sales_amounts)
total_sales
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[114], line 2
      1 # Find the total amount of sales for a whole day
----> 2 total_sales = prices.dot(sales_amounts)
      3 total_sales

ValueError: shapes (3,) and (5,3) not aligned: 3 (dim 0) != 5 (dim 0)
```

The shapes aren't aligned, we need the middle two numbers to be the same.

```
[115]: prices
```

```
[115]: array([10,  8, 12])
```

```
[116]: sales_amounts.T.shape
```

```
[116]: (3, 5)
```

```
[117]: # To make the middle numbers the same, we can transpose
total_sales = prices.dot(sales_amounts.T)
total_sales
```

```
[117]: array([240, 138, 458, 232, 142])
```

```
[118]: butter_prices.shape, weekly_sales.shape
```

```
[118]: ((1, 3), (5, 3))
```

```
[119]: daily_sales = butter_prices.dot(weekly_sales.T)
daily_sales
```

```
[119]:      Mon  Tues  Wed  Thurs  Fri
Price  240   138  458    232  142
```

```
[120]: # Need to transpose again
weekly_sales["Total"] = daily_sales.T
weekly_sales
```

	Almond butter	Peanut butter	Cashew butter	Total
Mon	12	15	0	240
Tues	3	3	7	138
Wed	9	19	18	458
Thurs	4	6	12	232
Fri	1	6	7	142

1.5.8 Comparison operators

Finding out if one array is larger, smaller or equal to another.

```
[121]: a1
```

```
[121]: array([1, 2, 3])
```

```
[122]: a2
```

```
[122]: array([[1. , 2. , 3.3],
             [4. , 5. , 6.5]])
```

```
[123]: a1 > a2
```

```
[123]: array([[False, False, False],
             [False, False, False]])
```

```
[124]: a1 >= a2
```

```
[124]: array([[ True,  True, False],
             [False, False, False]])
```

```
[125]: a1 > 5
```

```
[125]: array([False, False, False])
```

```
[126]: a1 == a1
```

```
[126]: array([ True,  True,  True])
```

```
[127]: a1 == a2
```

```
[127]: array([[ True,  True, False],
             [False, False, False]])
```

1.6 5. Sorting arrays

- `np.sort()` - sort values in a specified dimension of an array.
- `np.argsort()` - return the indices to sort the array on a given axis.
- `np.argmax()` - return the index/indices which gives the highest value(s) along an axis.
- `np.argmin()` - return the index/indices which gives the lowest value(s) along an axis.

```
[128]: random_array
```

```
[128]: array([[8, 7, 6],  
            [4, 2, 7],  
            [6, 0, 6],  
            [0, 8, 5],  
            [6, 2, 9]])
```

```
[129]: np.sort(random_array)
```

```
[129]: array([[6, 7, 8],  
            [2, 4, 7],  
            [0, 6, 6],  
            [0, 5, 8],  
            [2, 6, 9]])
```

```
[130]: np.argsort(random_array)
```

```
[130]: array([[2, 1, 0],  
            [1, 0, 2],  
            [1, 0, 2],  
            [0, 2, 1],  
            [1, 0, 2]])
```

```
[131]: a1
```

```
[131]: array([1, 2, 3])
```

```
[132]: # Return the indices that would sort an array  
np.argsort(a1)
```

```
[132]: array([0, 1, 2])
```

```
[133]: # No axis  
np.argmin(a1)
```

```
[133]: np.int64(0)
```

```
[134]: random_array
```

```
[134]: array([[8, 7, 6],  
            [4, 2, 7],  
            [6, 0, 6],  
            [0, 8, 5],  
            [6, 2, 9]])
```

```
[135]: # Down the vertical  
np.argmax(random_array, axis=1)
```

```
[135]: array([0, 2, 0, 1, 2])
```

```
[136]: # Across the horizontal  
np.argmin(random_array, axis=0)
```

```
[136]: array([3, 2, 3])
```

1.7 numbers6. Use case

Turning an image into a NumPy array.

Why?

Because computers can use the numbers in the NumPy array to find patterns in the image and in turn use those patterns to figure out what's in the image.

This is what happens in modern computer vision algorithms.

```
[19]: from IPython.display import Image, display  
image_path = 'tokyo.jpg'  
display(Image(filename=image_path))
```



```
[21]: from matplotlib.image import imread  
  
tokyo = imread('./tokyo.jpg')
```



```
print(type(tokyo))
```

```
<class 'numpy.ndarray'>
```

```
[22]: tokyo.shape
```

```
[22]: (334, 500, 3)
```

```
[23]: tokyo
```

```
[23]: array([[30, 53, 87],
            [30, 53, 87],
            [30, 53, 87],
            ...,
            [20, 21, 67],
            [20, 21, 67],
            [20, 21, 67]],

           [[30, 53, 87],
            [30, 53, 87],
            [30, 53, 87],
            ...,
            [21, 22, 68],
            [21, 22, 68],
            [21, 22, 68]],

           [[31, 54, 88],
            [31, 54, 88],
            [31, 54, 88],
            ...,
            [21, 22, 68],
            [21, 22, 68],
            [21, 22, 68]],

           ...,

           [[ 7,  7, 31],
            [ 9,  9, 33],
            [43, 43, 67],
            ...,
            [83, 49, 84],
            [73, 39, 74],
            [50, 16, 53]],

           [[62, 54, 77],
            [34, 26, 49],
            [45, 40, 63],
            ...,
```

```
[93, 56, 90],  
[72, 34, 71],  
[40, 4, 40]],  
  
[[77, 65, 87],  
[46, 37, 58],  
[52, 42, 67],  
...,  
[62, 22, 56],  
[44, 6, 43],  
[27, 0, 26]]], dtype=uint8)
```

```
[24]: image_path = 'supra.jpg'  
display(Image(filename=image_path))
```



```
[6]: car = imread("./supra.jpg")  
car.shape
```

```
[6]: (1726, 3068, 3)
```

```
[7]: car[:, :, :3].shape
```

```
[7]: (1726, 3068, 3)
```

2 Audio Manipulations With NumPy

The Best way to learn NumPy is to learn by doing!

Lets apply these concepts in a real world project.

```
[ ]: # Install librosa if you haven't
!pip install librosa
```

```
[2]: import librosa
from librosa.display import waveshow
from IPython.display import Audio
```


```
[4]: # load the audio into a numpy array
# y is the numpy array representing the waveform
# sr is the sampling rate of the waveform
y, sr = librosa.load(r"./Erika-German-Song-Sound-Effect.mp3")
```

```
[ ]: Audio(y,rate=sr) # allows you to play the audio withing jupyter
```

```
[ ]: y.ndim
```


```
[ ]: y.shape
```

```
[9]: waveshow(y,sr=sr);
```



```
[10]: # Reduce the amplitude  
y /= 50
```


```
[11]: waveshow(y,sr=sr);
```



```
[ ]: Audio(y,rate=sr)
```

```
[12]: # Reverse the audio  
y = y[::-1]
```

```
[14]: waveshow(y,sr=sr);
```



```
[ ]: Audio(y,rate=sr)
```

```
[ ]: # Lets add an echo effect to the soundtrack  
echo = np.zeros_like(y)  
echo = np.zeros(y.shape)  
echo.shape
```

```
[ ]: echo[int(sr*0.5):] = echo[:int(sr*0.5)]+0.7
```

```
[ ]: Audio(y+echo,rate=sr)
```

Libraries like NumPy can only be mastered by consistent usage and building. Practice writing code, making real world applications and just make projects for the sake of it!

3 Thank You!